

SMT-COMP 2022

17th International Satisfiability Modulo Theory Competition

Haniel Barbosa Jochen Hoenicke François Bobot

Universidade Federal de Minas Gerais, Brazil

Albert-Ludwigs-Universität Freiburg, Germany

CEA List, France

Aug 11, 2022

SMT-COMP

Annual competition for [SMT solvers](#)
on (a selection of) benchmarks from [SMT-LIB](#)

History

2005	first competition
2013	evaluation instead of competition
2014	since then hosted by StarExec

Goals:

- spur development of SMT solver implementations
- promote SMT solvers and their usage
- support the SMT-LIB project
 - to promote and develop the SMT-LIB format
 - model validation
 - proof checking
 - to collect relevant benchmarks
- engage and include new members

SMT Solvers and SMT-LIB

SMT Solver

- checks formulas in **SMT-LIB** format for **satisfiability modulo theories**

SMT-LIB is

- ① a **language** in which benchmarks are written
- ② a community effort to **collect benchmarks**

Non-incremental

391 363 instances (+9680)
with 1 query each
in 81 logics (+2).

Incremental

43 285 instances (+1)
with 33 998 935 queries (+141)
in 39 logics.

SMT Solvers and SMT-LIB

SMT Solver

- checks formulas in **SMT-LIB** format for **satisfiability modulo theories**

SMT-LIB is

- ① a **language** in which benchmarks are written
- ② a community effort to **collect benchmarks**

Non-incremental

391 363 instances (+9680)
with 1 query each
in 81 logics (+2).

Incremental

43 285 instances (+1)
with 33 998 935 queries (+141)
in 39 logics.

Selected Non-incremental

206 932 instances

Selected Incremental

22 300 instances

SMT Solvers and SMT-LIB

SMT Solver

- checks formulas in **SMT-LIB** format for **satisfiability modulo theories**

SMT-LIB is

- ① a **language** in which benchmarks are written
- ② a community effort to **collect benchmarks**

Non-incremental

391 363 instances (+9680)
with 1 query each
in 81 logics (+2).

Incremental

43 285 instances (+1)
with 33 998 935 queries (+141)
in 39 logics.


Selected Non-incremental

206 932 instances
We mistakenly ignored 9466 new
benchmarks

Selected Incremental

22 300 instances

Competition Overview



SMT-COMP Tracks (traditional)

Single Query (SQ) Track

- Determine satisfiability of one problem
- Solver answers sat/unsat/unknown

Unsat Core Track

- Find small unsatisfiable subset of input.
- Solver answers unsat + list of formulas.

Model Validation Track

- Find a model for a satisfiable problem.
- Solver answers sat + value for each non-logical symbol.

Incremental Track

- Solve many small problems interactively.
- Solver acks commands and answers sat/unsat for each check.

SMT-COMP Tracks (experimental)

Model Validation

- Division with quantifier-free floating-point logics
- Model validation with Dolmen (thanks to Gillaume Bury and François Bobot)

Cloud and Parallel Track (sponsored by AWS, led by Mike Whalen)

- Solve a large problem over the cloud (or a big computer)
 - 100 machines, 1600 cores, 6400 GB of memory (cloud)
 - 64 cores, 256 GB of memory (parallel)
- Solver answers sat/unsat/unknown

SMT-COMP Tracks (experimental)

Model Validation

- Division with quantifier-free floating-point logics
- Model validation with Dolmen (thanks to Gillaume Bury and François Bobot)

Cloud and Parallel Track (sponsored by AWS, led by Mike Whalen)

- Solve a large problem over the cloud (or a big computer)
 - 100 machines, 1600 cores, 6400 GB of memory (cloud)
 - 64 cores, 256 GB of memory (parallel)
- Solver answers sat/unsat/unknown

Proof Exhibition Track

- Solver submitted together with a checker for unsatisfiability proofs
- No predefined format or checker
- No ranking
- Qualitative assessment

SMT-COMP Tracks (experimental)

Model Validation

- Division with quantifier-free floating-point logics
- Model validation with Dolmen (thanks to Gillaume Bury and François Bobot)

Cloud and Parallel Track (sponsored by AWS, led by Mike Whalen)

- Solve a large problem over the cloud (or a big computer)
 - 100 machines, 1600 cores, 6400 GB of memory (cloud)
 - 64 cores, 256 GB of memory (parallel)
- Solver answers sat/unsat/unknown

Proof Exhibition Track

- Solver submitted together with a checker for unsatisfiability proofs
- No predefined format or checker
- No ranking
- Qualitative assessment

This year the sat/unsat results from sound solvers in SQ were used to include benchmarks on the MV, UC and PE tracks.

Tracks, Solvers, Divisions, and Benchmarks

Teams: 21 (+3)

Track	Solvers	Divisions	Benchmarks
Single Query	22(+3)	19(+1)	93 945
Incremental	8(+1)	17(+2)	22 300
Unsat Core	6(-1)	18(+1)	57 245
Model Validation	8(+1)	7(+ 1 exp.)	32 766
Proof Exhibition	4	18 exp.	57 245
Parallel	4(+1)	14 exp.	400
Cloud	4(-1)	14 exp.	400

Number in parenthesis shows changes from 2021

Participants

SMT-COMP 2022 participants rely on multiple reasoning frameworks:

- CDCL(T)
- mcSAT
- saturation
- automata
- finite domain
- CP
- local search
- besides wrappers extending the scope of existing solvers

Six new solvers participated:

- NRA-LS (Liu et al.)
- OSTRICH (Chen et al.)
- Yices-ismt (Jia et al.)
- Z3++ (Cai et al.)
- solsmt (Reitwiessner and Soos)

Solver Presentation



Tracks/Divisions

Single Query:	$\text{QF_}\{A, BV, FP, FPLRA, UF\}^+, \{A, BV, FP, FPLRA, UF\}^+$
Incremental:	$\text{QF_}\{A, BV, FP, FPLRA, UF\}^+, \{A, BV, FP, FPLRA, UF\}^+$
Unsat Core:	$\text{QF_}\{A, BV, FP, FPLRA, UF\}^+$
Model Validation:	$\text{QF_}BV, \text{QF_}BVFP, \text{QF_}BVFPLRA, \text{QF_}FP, \text{QF_}FPLRA, \text{QF_}UFBV, \text{QF_}UFPF$

Highlights

- Quantifiers support for all supported theories (**new**)
- Sequential combination of bit-blasting and propagation based local search for $\text{QF_}BV$ (SQ, MV)
- Uses CaDiCaL version 1.5.2 as SAT backend

<https://bitwuzla.github.io>

COLIBRI

Bruno Marre, François Bobot

CP solver:

- Domains (Intervals, modulo, FP and Bitvector representation)
 - Local propagators for every constraints
 - Constraint rewriting
 - Global propagators (Simplex,...)
 - Decision by domain splitting or value
- FP, Modulo arithmetic, reals, BV
 - $\leq 25s$ QF_FP
 - Fuzzing with non-standard operators?
 - computer division
 - Encoding!

CVC5 at the SMT Competition 2022

H. Barbosa, C. Barrett, M. Brain, G. Kremer, H. Lachnitt, A. Mohamed, M. Mohamed, A. Niemetz,
A. Nötzli, A. Ozdemir, M. Preiner, A. Reynolds, Y. Sheng, C. Tinelli, Y. Zohar

Versatile and Industrial-Strength

- ▶ Support for all standardized SMT-LIB theories, user-friendly API
- ▶ Features beyond SMT solving (synthesis, proofs, ...)
- ▶ This year: Focus on robustness, features

Configurations

CVC5 entered all divisions in all tracks.

- ▶ Single query track: Sequential portfolio
- ▶ Unsat-core track: Based on proof module or assumptions in the SAT solver

Proof Exhibition Track

- ▶ Two configurations: Internal proof checker and LFSC (CVC5-lfsc)
- ▶ Default
 - ▶ Uses the internal proof format: directed acyclic graphs of proof rule applications
 - ▶ Proofs are checked during construction
- ▶ CVC5-lfsc
 - ▶ Uses CVC5's LFSC back end
 - ▶ Proof and proof signatures in LFSC
 - ▶ Proof checker ensures that proof is well-formed w.r.t. the signature

Components

- ▶ Divide-and-conquer algorithm
- ▶ CVC5 as the base solver
- ▶ A splitter based on CVC5
- ▶ An MPI-based architecture for scheduling

Splitter

- ▶ Uses existing infrastructure and smarts of CVC5
- ▶ Intercepts calls to theory solvers after configurable number of checks
- ▶ Collects subset of literals l_1, \dots, l_m from the current decision trail
- ▶ Blocks $\neg(l_1 \wedge \dots \wedge l_m)$
- ▶ Generates n -th cube $l_1 \wedge \dots \wedge l_m \wedge \neg C_1 \wedge \dots \wedge C_{n-1}$
- ▶ If less than two partitions are made, tries with other parameters

NRA-LS: Applying Local Search on Non-linear Real Arithmetic

Minghao Liu, Fuqi Jia, Rui Han, Yu Zhang, Pei Huang, Feifei Ma, Jian Zhang



- New entrant to SMT-COMP, participating in QF_NRA single query track this year

Highlights


- A local search solver that wraps CVC5-1.0.0
- To solve satisfiable instances with high-order polynomials

Approach

- Initial model generation
- Iterative satisfiability testing of sub-problems
- Time slots assignment

Resources

- System description and project are available at
<https://github.com/minghao-liu/NRA-LS>



OpenSMT 2022

Efficient, interpolating solver for linear arithmetic, uninterpreted functions, and arrays

Written in C++17, used by the Horn solver Golem

New logics:

Combination logics QF_UFLIA and QF_UFLRA
(model-based theory combination)
Arrays QF_AX (from SMTInterpol)

Performance improvements

Arithmetic: Cuts-from-proofs,
memory efficiency
SAT solver: phase saving,
glucose-style learned clauses
management

Proof track

Custom theory-specific trail format
with drat-based SAT proofs



Parallel and cloud solver (SMTS)

Based on the partition-tree approach
Search-space partitioning
Clause sharing between arbitrary partitions
Major rewrite from last year
Two versions: portfolio and cube-and-conquer


Straight-Line Formulas

- ▶ $S ::= x := f(\bar{x}) \mid \mathbf{assert}(R(\bar{x})) \mid S; S$
- ▶ Soundness + completeness guarantee for SL formulas
- ▶ Complex string functions, including replaceall and transducers

Backwards-propagation

- ▶ Main algorithm for SL formulas
- ▶ $z = x \circ y; \mathbf{assert}(z \in L)$
- ▶ $\mathbf{assert}(x \in L_1); \mathbf{assert}(y \in L_2)$

Extensions



Interpolating SMT solver

- based on CDCL(T)
- for Arrays, Uninterpreted Functions, Linear Integer and Real Arithmetic
 - plus `div` and `mod` with constants, and Data Types
- supports quantifiers
- produces models, proofs, and unsat cores
- computes sequence and tree interpolants

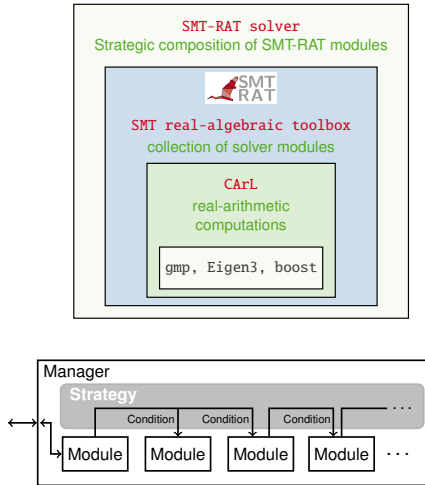
SMTInterpol at SMT-COMP 2022

- with **detailed proof production**
- **proof check** in single query/unsat core
- **quantified data type** logics

Try it in your browser:



<https://tinyurl.com/smtinterpol>



CArL library: basic arithmetic datatypes and computations

[Sapientia'18, NFM'11, CAI'11]

Basic modules

Preprocessing/simplifying modules

SAT solver

CNF converter

Non-algebraic decision procedures

Bit-blasting

Equalities and uninterpreted functions

Bit-vectors

Interval constraint propagation

Pseudo-Boolean formulas

Algebraic decision procedures

Gauß+Fourier-Motzkin

Simplex [ISSAC'21]

Linearization

Gröbner bases [CAI'13]

MCSAT (FM,VS,CAD) [2xSC²'19]

Subtropical satisfiability

Virtual substitution [SC²'17, FCT'11; PhD Corzilius]

Cylindrical algebraic decomposition [JLAMP'21, SYNASC'21, SC²'21, CADE-24, JSC'20, SC²'20, SC²'17; PhDs Loup, Kremer, Nalbach]

Generalized branch-and-bound [CASC'16]

Cube tests

SoISMT

tiny SMT solver inside the Solidity compiler

Christian Reitwiessner, Mate Soos
Ethereum Foundation



Solidity: Most widely used language for Ethereum smart contracts

not to be confused with **SoICMC**: interface to other SMT solvers for program verification

SoISMT: tiny integrated SMT solver used in optimizer

- remove redundant overflow checks, determine non-overlapping memory access, etc
- any bug in SMT solver can lead to bug in program
- needs to be fully deterministic and platform-independent for reproducibility
- implements QF_LRA using CDCL(T)
- inspired by MiniSat1.14 plus Dutertre-deMoura for LRA
- written in ~3k lines of C++

plan to add proof generation and checking for SAT and theory

- STP is an eager bit-blasting solver for QF_BV, and also QF_ABV but without extensionality.
- Our focus is on analysis at the bit-vector level, for example with unsigned-intervals, then applying sharing-aware rewrites. Because of this we go well on inefficiently-encoded problems.
- Development is hosted on Github.



Bruno Andreotti¹, Haniel Barbosa¹, Pascal Fontaine^{2,3}, *Hans-Jörg Schurr*³

¹Department of Computer Science, Universidade Federal de Minas Gerais (UFMG)

²CNRS, Inria, and the University of Lorraine, Nancy, France

³Université de Liège, Belgium

- ▶ Focus on proofs
 - ▶ Participates in the proof exhibition
 - ▶ Uses the Alethe format
 - ▶ Proofs are checked with an high-performance proof checker written in Rust
 - ▶ Many small fixes and clarifications

Yices 2 in SMTCOMP 2021

Yices 2

- Supports linear and non-linear arithmetic, arrays, UF, bitvectors
- Supports incremental solving and unsat cores
- Includes two types of solvers: classic CDCL(T) + MCSAT
- <https://github.com/SRI-CSL/yices2>
- <https://yices.csl.sri.com>

New in 2021

- Quantifier reasoning: model-based quantifier instantiation + E-graph matching (thanks to Aman Goel)
- MCSAT extensions
 - Solving modulo a model
 - Interpolant for MCSAT-supported theories

Fuqi Jia, Rui Han, Minghao Liu, Cunjing Ge, Pei Huang, Feifei Ma, Jian Zhang.

- New entrant to SMT-COMP, participating in QF_NIA single query track this year.

- A wrapper solver: Yices2 + ismt.

- ismt: Preprocessor + “The Three Musketeers”.

- **Decider**: decide abstraction assignment \rightarrow Interval Assignment.
- **Searcher**: search in the given space \rightarrow Bit-Blasting.
- **Resolver**: resolve conflict from failure \rightarrow Failed Interval Lemma ψ .

Dependencies

- *Yices 2.6.2*
- *Libpoly v0.1,11*
- *CaDiCal 1.5.2*

- Yices-ismt: $\text{Yices2}(\phi) \rightarrow \text{ismt}(\phi) \rightarrow \text{Yices}(\phi \wedge \psi)$.

- ψ rule out failed space.

- ◆ <https://github.com/MRVAPOR/Yices-ismt>

YicesQS, an extension of Yices2 for quantifiers (SMT-comp 2022)


Stéphane Graham-Lengrand

<https://github.com/disteph/yicesQS>

YicesQS implements a 2-player game (\forall player vs \exists player) playing on a quantified input formula F . Our recursive generalization of counter-example-guided quantifier instantiation (CEGQI) produces a quantifier-free satisfiable under-approximation of F or a quantifier-free unsatisfiable over-approximation of F .

2022: YicesQS entered logics NRA , NIA ^{NEW}, LRA ^{NEW}, LIA ^{NEW} and BV , & generally targets complete theories with procedures for answering 3 types of quantifier-free queries:

- *Satisfiability modulo assignment / modulo a model* (here relying on MCSAT)
- *Model generalization* (here using invertibility conditions for BV, CAD projections + ϵ -terms for algebraic reals^{NEW} for arithmetic)
- *Model interpolation* (here again relying on MCSAT, incl. CAD for arithmetic)



YicesQS is written in OCaml, using Yices2 as a library via its OCaml bindings.

<https://github.com/SRI-CSL/yices2>

https://github.com/SRI-CSL/yices2_ocaml_bindings





<https://z3-plus-plus.github.io/>

Shaowei Cai, Bohan Li, Jinkun Lin, Zhonghan Wang, Bohua Zhan, Xindi Zhang, Mengyu Zhao

State Key Laboratory of Computer Science

Institute of Software, Chinese Academy of Sciences, Beijing, China



QF_BV:

Off-the-shelf SAT solvers;



Word-level and bit-level rewriting rules

QF_LIA, QF_IDL, QF_NIA:

A local search solver dedicated for integer arithmetic logic



QF_NRA:

A feasible region consistency checker;

Sample-cell projection in NLSAT



Other participants

- Q3B
- SMT-RAT
- SMTS
- UltimateEliminator+MathSAT
- Vampire
- Z3str4

Non-Competitive Solvers

Submitted by organisers

- z3-4.8.17
- MathSAT 5.6.8
- Best solvers, per division, from previous years (23 Solvers)

Submitted by participants

- Fixed solvers (OpenSMT, STP, Yices-ismt, Z3++,smtinterpol)

Scoring

Computing scores:

- **Single Query/Parallel/Cloud**: number of solved **instances**
- **Incremental**: number of solved **queries**
- **Unsat Core**: number of top-level assertions **removed**
- **Model Validation**: number of solved instances with correct **models**

Error scores:

- **All Tracks**: given for sat reply for unsat instance, or vice versa
- **Unsat Core**: given if returned core is satisfiable.
- **Model Validation**: given if given model evaluates formula to **false**

Error scores are draconian.

Score and Ranking

In each track we collect different scores:

- **Sequential score** (SQ, UC, MV): all time limits apply to cpu time
- **Parallel score** (all): all time limits apply to wallclock time
- **SAT score** (SQ): parallel score for **satisfiable** instances
- **UNSAT score** (SQ): parallel score for **unsatisfiable** instances
- **24s** (SQ): parallel score with time limit of **24s**

Division ranking (for each score)

- For each division, one winner is declared

Two competition-wide rankings (for each score)

- **Biggest lead**: division winner with most score difference to second place
- **Largest contribution**: improvement each solver provided to a virtual best solver

Division Winners

Division Winners

Single Query, sequential score

- **Bitwuzla**: FPArith, QF_Bitvec, QF_Equality+Bitvec, QF_FPArith
- **cvc5**: Arith, Bitvec, Equality, Equality+LinearArith, Equality+MachineArith, Equality+NonLinearArith, QF_Datatypes, QF_Equality+NonLinearArith, QF_NonLinearIntArith, QF_NonLinearRealArith, QF_Strings
- **OpenSMT**: QF_LinearIntArith
- **SMTInterpol**: QF_Equality+LinearArith
- **Yices2**: QF_Equality, QF_LinearRealArith

Division Winners

Single Query, sequential score

- **Bitwuzla**: FPArith, QF_Bitvec, QF_Equality+Bitvec, QF_FPArith
- **cvc5**: Arith, Bitvec, Equality, Equality+LinearArith, Equality+MachineArith, Equality+NonLinearArith, QF_Datatypes, QF_Equality+NonLinearArith, QF_NonLinearIntArith, QF_NonLinearRealArith, QF_Strings
- **OpenSMT**: QF_LinearIntArith
- **SMTInterpol**: QF_Equality+LinearArith
- **Yices2**: QF_Equality, QF_LinearRealArith

Unsat Core

- **cvc5**: Arith, Bitvec, Equality+LinearArith, Equality+MachineArith, Equality+NonLinearArith, Equality(Seq), QF_Datatypes
- **Vampire**: Equality(Par)
- **Bitwuzla**: QF_EqualityBitvec, QF_FPArith
- **Yices2**: QF_Bitvec, QF_Equality+LinearArith, QF_Equality, QF_LinearIntArith, QF_LinearRealArith
- **SMTInterpol**: QF_Equality+NonLinearArith
- **UltimateEliminator+MathSAT**: FPArith

Division Winners

Incremental

- **cvc5**: Arith, Bitvec, Equality+LinearArith, EqualityNonLinearArith, Equality
- **UltimateEliminator+MathSAT**: Equality+MachineArith
- **Bitwuzla**: FPArith, QF_FPArith
- **Yices2**: QF_Bitvec, QF_Equality+Bitvec, QF_Equality, QF_LinearIntArith
- **SMTInterpol**: QF_Equality+LinearArith, QF_Equality+NonLinearArith, QF_NonLinearIntArith
- **OpenSMT**: QF_LinearRealArith

Division Winners

Incremental

- **cvc5**: Arith, Bitvec, Equality+LinearArith, EqualityNonLinearArith, Equality
- **UltimateEliminator+MathSAT**: Equality+MachineArith
- **Bitwuzla**: FPArith, QF_FPArith
- **Yices2**: QF_Bitvec, QF_Equality+Bitvec, QF_Equality, QF_LinearIntArith
- **SMTInterpol**: QF_Equality+LinearArith, QF_Equality+NonLinearArith, QF_NonLinearIntArith
- **OpenSMT**: QF_LinearRealArith

Model Validation (competitive only)

- **Bitwuzla**: QF_Bitvec, QF_Equality+Bitvec,
- **smtinterpol**: QF_Equality+LinearArith
- **Yices2**: QF_Equality
- **Z3++**: QF_LinearIntArith
- **OpenSMT**: QF_LinearRealArith
- **experimental**: QF_FPArith : Bitwuzla, cvc5

Largest contribution

	1st Place		2nd Place		3rd Place	
Single Query						
seq	cvc5	(Eq+MA)	YicesQS	(Arith)	Bitwuzla	(FPArith)
par	cvc5	(Eq+MA)	YicesQS	(Arith)	Bitwuzla	(FPArith)
sat	cvc5	(Eq+LA)	YicesQS	(Arith)	Bitwuzla	(FPArith)
unsat	cvc5	(Eq+MA)	Z3++	(QF_NonLIA)	OSTRICH	(QF_Strings)
24	Vampire	(Eq+NA)	Vampire	(Equality)	Yices2	(QF_LinIA)

Largest contribution

	1st Place	2nd Place	3rd Place
Single Query			
seq	cvc5 (Eq+MA)	YicesQS (Arith)	Bitwuzla (FPArith)
par	cvc5 (Eq+MA)	YicesQS (Arith)	Bitwuzla (FPArith)
sat	cvc5 (Eq+LA)	YicesQS (Arith)	Bitwuzla (FPArith)
unsat	cvc5 (Eq+MA)	Z3++ (QF_NonLIA)	OSTRICH (QF_Strings)
24	Vampire (Eq+NA)	Vampire (Equality)	Yices2 (QF_LinIA)
Incremental			
par	cvc5 (Eq+NA)	Yices2 (QF_Eq+LA)	SMTInterpol(QF_Eq+NA)

Largest contribution

	1st Place	2nd Place	3rd Place
Single Query			
seq	cvc5 (Eq+MA)	YicesQS (Arith)	Bitwuzla (FPArith)
par	cvc5 (Eq+MA)	YicesQS (Arith)	Bitwuzla (FPArith)
sat	cvc5 (Eq+LA)	YicesQS (Arith)	Bitwuzla (FPArith)
unsat	cvc5 (Eq+MA)	Z3++ (QF_NonLIA)	OSTRICH (QF_Strings)
24	Vampire (Eq+NA)	Vampire (Equality)	Yices2 (QF_LinIA)
Incremental			
par	cvc5 (Eq+NA)	Yices2 (QF_Eq+LA)	SMTInterpol(QF_Eq+NA)
Unsat Core			
seq	cvc5 (Eq+LA)	Bitwuzla (QF_Eq+Bitvec)	
par	cvc5 (Eq+NA)	Bitwuzla (QF_Eq+Bitvec)	

Largest contribution

	1st Place	2nd Place	3rd Place
Single Query			
seq	cvc5 (Eq+MA)	YicesQS (Arith)	Bitwuzla (FPArith)
par	cvc5 (Eq+MA)	YicesQS (Arith)	Bitwuzla (FPArith)
sat	cvc5 (Eq+LA)	YicesQS (Arith)	Bitwuzla (FPArith)
unsat	cvc5 (Eq+MA)	Z3++ (QF_NonLIA)	OSTRICH (QF_Strings)
24	Vampire (Eq+NA)	Vampire (Equality)	Yices2 (QF_LinIA)
Incremental			
par	cvc5 (Eq+NA)	Yices2 (QF_Eq+LA)	SMTInterpol (QF_Eq+NA)
Unsat Core			
seq	cvc5 (Eq+LA)	Bitwuzla (QF_Eq+Bitvec)	
par	cvc5 (Eq+NA)	Bitwuzla (QF_Eq+Bitvec)	
Model Validation			
seq	Z3++ (QF_LinIA)	Bitwuzla (QF_Bitvec)	smtinterpol (QF_Eq+LIA)
par	Z3++ (QF_LinIA)	Bitwuzla (QF_Bitvec)	smtinterpol (QF_Eq+LIA)

Biggest Lead

	1st Place		2nd Place		3rd Place	
Single Query						
seq	cvc5	(QF_DT)	Bitwuzla	(FPArith)	Yices2	(QF_LinRA)
par	cvc5	(QF_DT)	Bitwuzla	(FPArith)	Yices2	(QF_LinRA)
sat	cvc5	(QF_DT)	Bitwuzla	(FPArith)	Z3++	(QF_NonRA)
unsat	cvc5	(QF_DT)	Z3++	(QF_NonIA)	Bitwuzla	(FPArith)
24	cvc5	(Eq+MA)	smtinterpol	(QF_DT)	Vampire	(Equality)

Biggest Lead

	1st Place	2nd Place	3rd Place
Single Query			
seq	cvc5 (QF_DT)	Bitwuzla (FPArith)	Yices2 (QF_LinRA)
par	cvc5 (QF_DT)	Bitwuzla (FPArith)	Yices2 (QF_LinRA)
sat	cvc5 (QF_DT)	Bitwuzla (FPArith)	Z3++ (QF_NonRA)
unsat	cvc5 (QF_DT)	Z3++ (QF_NonIA)	Bitwuzla (FPArith)
24	cvc5 (Eq+MA)	smtinterpol(QF_DT)	Vampire (Equality)
Incremental			
par	SMTInterpol(QF_NIA)	Yices2 (QF_LinIA)	cvc5 (Eq+NonLA)

Biggest Lead

	1st Place	2nd Place	3rd Place
Single Query			
seq	cvc5 (QF_DT)	Bitwuzla (FPArith)	Yices2 (QF_LinRA)
par	cvc5 (QF_DT)	Bitwuzla (FPArith)	Yices2 (QF_LinRA)
sat	cvc5 (QF_DT)	Bitwuzla (FPArith)	Z3++ (QF_NonRA)
unsat	cvc5 (QF_DT)	Z3++ (QF_NonIA)	Bitwuzla (FPArith)
24	cvc5 (Eq+MA)	smtinterpol(QF_DT)	Vampire (Equality)
Incremental			
par	SMTInterpol(QF_NIA)	Yices2 (QF_LinIA)	cvc5 (Eq+NonLA)
Unsat Core			
seq	cvc5 (Eq+MA)	Yices2 (QF_Eq+LA)	smtinterpol(QF_Eq+NA)
par	cvc5 (Eq+MA)	Yices2 (QF_Eq+LA)	smtinterpol(QF_Eq+NA)

Biggest Lead

	1st Place		2nd Place		3rd Place	
Single Query						
seq	cvc5	(QF_DT)	Bitwuzla	(FPArith)	Yices2	(QF_LinRA)
par	cvc5	(QF_DT)	Bitwuzla	(FPArith)	Yices2	(QF_LinRA)
sat	cvc5	(QF_DT)	Bitwuzla	(FPArith)	Z3++	(QF.NonRA)
unsat	cvc5	(QF_DT)	Z3++	(QF.NonIA)	Bitwuzla	(FPArith)
24	cvc5	(Eq+MA)	smtinterpol	(QF_DT)	Vampire	(Equality)
Incremental						
par	SMTInterpol(QF_NIA)		Yices2	(QF_LinIA)	cvc5	(Eq+NonLA)
Unsat Core						
seq	cvc5	(Eq+MA)	Yices2	(QF.Eq+LA)	smtinterpol	(QF.Eq+NA)
par	cvc5	(Eq+MA)	Yices2	(QF.Eq+LA)	smtinterpol	(QF.Eq+NA)
Model Validation						
seq	Z3++	(QF_LinIA)	smtinterpol	(QF_LinA)	OpenSMT	(QF_LinRA)
par	Z3++	(QF_LinIA)	smtinterpol	(QF_LinA)	OpenSMT	(QF_LinRA)

Checking Disagreements

- 111 451 instances of 391 363 have no status
- Only 18 benchmarks with disagreements (AUFDTLIA, QF_NIA)
- We manually resolved the disagreements
 - Authors confirmed solver unsoundness
 - Sound solvers agreed on result
- We had only three solvers with soundness issues: SMTIinterpol, Z3++, yices-ismt
 - Down from 10 last year

Checking Disagreements

- 111 451 instances of 391 363 have no status
- Only 18 benchmarks with disagreements (AUFDTLIA, QF_NIA)
- We manually resolved the disagreements
 - Authors confirmed solver unsoundness
 - Sound solvers agreed on result
- We had only three solvers with soundness issues: SMTIinterpol, Z3++, yices-ismt
 - Down from 10 last year (but note that this may be because of our issue with ignoring most of the new benchmarks in the selection)

Plans for SMT-COMP 2022

- Move to Dolmen as single model validator
- Dolmen worked well in QF_FP (minor issue with empty output, simple to fix)
- pysmt timed out in two cases (locally with more time one of them succeeded)
- pysmt could not understand the model in 9 cases
- Dolmen is easy to extend to new theories
- Easy to integrate with correct-by-construction code extracted from Coq
- Should be more efficient for handling function call

Plans for SMT-COMP 2022

- Proof *validation* track
- Hopefully proof exhibition this year will help
- We have to analyze the data still
 - Job only finished last week (took 18 days to run)
 - 830gb

SMT-COMP organizing committee

Three people organize the SMT-COMP. In 2022:

- Haniel Barbosa
- Jochen Hoenicke
- François Bobot

Both Haniel and Jochen have been organizers for three-years. One of them should be replaced.

We need a successor for next year's competition. Contact us if you would like to volunteer!

Acknowledgements

- [Andrea Micheli](#): pysmt
- [Guillaume Bury](#): Model Validator
- [Clark Barrett](#), [Pascal Fontaine](#), [Aina Niemetz](#), [Mathias Preiner](#), [Hans-Jörg Schurr](#): SMT-LIB benchmarks
- [Aaron Stump](#): StarExec support
- [Mike Whalen and team](#): Cloud/Parallel Track



Benchmark contributors

In 2022 [new benchmarks](#) were contributed by:

- Alex Coffin
- Alex Ozdemir
- Ali Uncu, James Davenport and Matthew England
- Bohan Li
- Elizabeth Polgreen
- Fuqi Jia
- Johann-Tobias Aaron and Raphael Schäg
- Matthew England and Miguel Del Rio Almajano
- Nicolas Amat
- Yannick Moy
- Yoni Zohar

Thanks

to all participants

Thanks

to all participants

and to you for listening